

[

Rich Context Competition Phase 2 (RCC-05)

Wolfgang Otto^{1*}, Andrea Zielinski¹, Behnam Ghavimi¹, Dimitar Dimitrov¹,
Narges Tavakolpoursaleh¹, Karam Abdulahhad¹, Katarina Boland¹, Stafan Dietze¹

Abstract

This document describes the approach, results, and software submitted by team RCC-05 to the Rich Context Competition (RCC). The team consists of members of the department *Knowledge Technologies in the Social Sciences* of *GESIS - Leibniz Institute for the Social Sciences*. The goal of the RCC is to automatically discover and link research datasets, methods, and fields in social science publications.

Keywords

Dataset Mention Extraction — Dataset Linking — Research Method Extraction — Research Field Classification

¹ Email: firstname.lastname@gesis.org. Affiliation: *Knowledge Technologies in the Social Sciences, GESIS – Leibniz Institute for the Social Sciences, Cologne, Germany*

*Corresponding author: wolfgang.otto@gesis.org

Contents

1	GESIS	1
1.1	Introduction	1
1.2	Approach, data and pre-processing	2
1.3	Dataset extraction	4
1.4	Research method extraction	5
1.5	Research field classification	9
1.6	Discussion	10
1.7	Conclusion	11
1.8	Acknowledgments	11
	Acknowledgments	11
1.9	References	11
	References	11

]

1. GESIS

Authors: Wolfgang Otto, Andrea Zielinski, Behnam Ghavimi, Dimitar Dimitrov, Narges Tavakolpoursaleh, Karam Abdulahhad, Katarina Boland, Stafan Dietze

Affiliation: *Knowledge Technologies in the Social Sciences, GESIS – Leibniz Institute for the Social Sciences, Cologne, Germany*

Corresponding author: wolfgang.otto@gesis.org

1.1 Introduction

GESIS - the Leibniz Institute for the Social Sciences (GESIS)¹ is the largest European research and infrastructure provider for the social sciences and offers research, data, services and infrastructures supporting all stages of the scientific process. The GESIS department *Knowledge Technologies for the Social Sciences (WTS)*² is responsible for developing all digital services

¹ <https://www.gesis.org/en/institute>

² <https://www.gesis.org/en/institute/departments/knowledge-technologies-for-the-social-sciences/>

and research data infrastructures at GESIS and aims at providing integrated access to social sciences data and services. Next to traditional social sciences research data, such as surveys and census data, an emerging focus is to build data infrastructures able to exploit novel forms of social sciences research data, such as large Web crawls and archives.

Research at WTS³ addresses areas such as information retrieval, information extraction & Natural Language Processing (NLP), semantic technologies and human computer interaction and aims at ensuring access and use of social sciences research data along the FAIR principles, for instance, through interlinking of research data, established vocabularies and knowledge graphs and by facilitating semantic search across distinct platforms and datasets. Due to the increasing importance of Web- and W3C standards as well as Web-based research data platforms, in addition to traditional research data portals, findability and interoperability of research data across the Web constitutes one current challenge. In the context of Web-scale reuse of social sciences resources, the extraction of structured data about scholarly entities such as datasets and methods from unstructured and semi-structured text, as found in scientific publications or resource metadata, is crucial in order to be able to uniquely identify social sciences resources and to understand their inherent relations.

Prior works at WTS/GESIS addressing such challenges apply NLP and machine learning techniques to, for instance, extract and disambiguate mentions of datasets⁴ [1, 2]), authors [3, 4] or software tools [5] from scientific publications or to extract and fuse of scholarly data from large-scale Web crawls [6, 7]. Resulting pipelines and data are used to empower scholarly search engines such as the *GESIS-wide search*⁵ [8] which provides federated search for scholarly resources (datasets, publications etc) across a range of GESIS information systems or the *GESIS DataSearch* platform⁶ [9], which enables search across a vast number of social sciences research datasets mined from the Web.

Given the strong overlap of our research and development profile with the recent initiatives of the Coleridge Initiative to evolve this research field through the Rich Context Competition (RCC)⁷, we are enthusiastic about having participated in the RCC2018 and are looking forward to continuing this collaboration towards providing sound frameworks and tools which automate the process of interlinking and retrieving scientific resources.

The central tasks in the RCC are the extraction and disambiguation of mentions of datasets and research methods as well as the classification of scholarly articles into a discrete set of research fields. After the first phase, each team received feedback from the organizers of the RCC consisting of a quantitative and qualitative evaluation. Whereas quantitative results of our initial contribution throughout phase one have shown significant room for improvement, the qualitative assessment, conducted by four judges on a sample of ten documents, underlined the potential of our approach.

Having been among the four shortlisted submissions for the second phase of the RCC, this chapter describes our approaches, techniques, and additional data used to address all three tasks. As described in the following subsections, we decided to follow a module-based approach where either individual modules or the entire pipeline can be reused. The remaining chapter is organised as follows. The following Section 1.2 provides an overview of our approach, used background data and preprocessing steps, whereas Sections 1.3, 1.4 and 1.5 describe our approaches in more detail, including results towards each of the tasks. Finally, we discuss our results in Section 1.6 and provide an overview of future work in Section 1.7.

1.2 Approach, data and pre-processing

This section describes the external data sources we used as well as our pre-processing steps.

1.2.1 Approach overview and initial evaluation feedback

The central tasks in the RCC are the extraction of dataset mentions from text. Even so, we considered the discovery of research methods and research fields important. To this end, we decided to follow a module-based approach. Users could choose to use each specific module solely or as parts of a data processing pipeline. Figure 1 shows an overview of modules developed and their dependencies. Here, the upper three modules (which are in gray) describe the pre-processing steps (cf. Section 1.2.3). The lower four modules (blue) are used to generate the output in a predefined format as specified by the competition.

The pre-processing step consists of extracting metadata and raw text from PDF documents. The output of this step is then used by the software modules responsible for tackling the individual sub-tasks. These sub-tasks are to discover research datasets (cf. Section 1.3), methods (cf. Section 1.4) and fields (cf. Section 1.5). First, a Named Entity Recognition module is used to find dataset mentions. This module used a supervised approach trained on a weakly labeled corpus. In the next step, we combine all recognized mentions for each publication and compare these mentions to the metadata from the list of datasets given by the competition. For this linking step the mentions and year information located in the same sentence are used. The corresponding sentence and extracted information are saved for debugging and potential usage in future pipeline components. The task of identifying research methods is done by an adjusted NER module trained on a corpus of scientific publications. For identifying

³ <https://www.gesis.org/en/research/applied-computer-science/labs/wts-research-labs>

⁴ <https://www.gesis.org/en/research/external-funding-projects/archive/infolis-i-and-ii>

⁵ search.gesis.org

⁶ <https://datasearch.gesis.org/>

⁷ <https://coleridgeinitiative.org/richcontextcompetition>

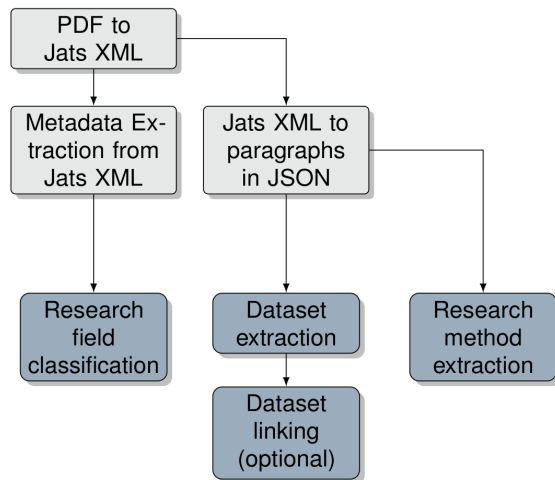


Figure 1. An overview of the individual software modules described in this document and their dependencies. 1- Gray: Our pre-processing pipeline. 2- Blue: three main tasks of the RCC.

research fields, we trained a classifier on openly available abstracts and metadata from the domain of social sciences crawled from the Social Science Open Access Repository⁸ (SSOAR). We tried different classifiers and selected the best performing one, a classifier based on fasttext⁹, i.e. a neural net based approach with a high performance[10].

After the first phase, each team received feedback from the organizers of the RCC. The feedback is two folds, a quantitative and qualitative evaluation. Unfortunately, the quantitative assessment showed our algorithms did not perform well regarding precision and recall metrics. In contrast to this, our approach has been found convincing regarding the quality of results. The qualitative feedback was from a random sample of ten documents given to four judges. Judges are then asked to manually extract dataset mentions and calculate the overlap between their dataset extractions and the output of our algorithm. Other factors that judges took into consideration are specificity, uniqueness, and multiple occurrences of dataset mentions. As for the extraction of research methods and fields, no ground truth has been provided; these tasks were evaluated against the judges' expert knowledge. Similarly to the extraction of dataset mentions, specificity and uniqueness have been considered for these two tasks. The feedback our team received acknowledged the fact that no ground truth has been provided and our efforts regarding the extraction of research methods and fields.

1.2.2 External data sources

For developing our algorithms, we utilized two external data sources. For the discovery of research methods and fields, we resort to data from the Social Science Open Access Repository¹⁰ (SSOAR). GESIS – Leibniz Institute for the Social Sciences maintains SSOAR by collecting and archiving literature of relevance to the social sciences.

In SSOAR, full texts are indexed using controlled social science vocabulary (Thesaurus¹¹, Classification¹²) and are assigned rich metadata. SSOAR offers documents in various languages. The corpus of English language publications that can be used for purposes of the competition consists of a total of 13,175 documents. All SSOAR documents can be accessed through the OAI-PMH¹³ interface.

Another external source, that we used for the discovery of research methods is the ACL Anthology Reference Corpus [11]. ACL ARC is a corpus of scholarly publications about computational linguistics. The corpus consists of a total of 22,878 articles.

1.2.3 Pre-processing

Although the organizers of the RCC offered plain texts for the publication, we decided to build our own pre-processing pipeline. The extraction of text from PDF files is still an error prone process. To handle de-hyphenation and paragraph segmentation during extraction time and benefit from automatic metadata extraction (i.e. title, author, abstracts and references) we decided to

⁸ <https://www.ssoar.info>

⁹ <https://fasttext.cc/>

¹⁰ <https://www.gesis.org/ssoar/home>

¹¹ <https://www.gesis.org/en/services/research/tools/thesaurus-for-the-social-sciences>

¹² <https://www.gesis.org/angebot/recherchieren/tools-zur-recherche/klassifikation-sozialwissenschaften> (in German)

¹³ <http://www.openarchives.org>

use a third party extraction tool. The Cermin Extraction Tool¹⁴[12] transforms the files into XML documents using the Journal Article Tag Suite¹⁵(Jats). For the competition we identified two interesting elements of the Jats XML format, i.e., <front> and <body>. The <front> element contains the metadata of the publication, whereas the <body> contains the main textual and graphic content of the publication. As a last step of the pre-processing, we removed all linebreaks from the publication. The output of this step is a list of metadata fields and values, as shown in Table 1 for each publication paragraph.

Table 1. Example preprocessing output for a paragraph in a given publication.

Example Text Field Data	
publication_id	12744
label	paragraph_text
text	A careful reading of text, word for word, was ...
section_title	Data Analysis
annotations	[{'start': 270, 'end': 295, 'type': 'bibref', ...
section_nr	[3, 2]
text_field_nr	31
para_in_section	1

1.3 Dataset extraction

1.3.1 Task description

In the scientific literature, datasets are cited to reference, for example, the data on which an analysis is performed or on which a particular result or claim is based. In this competition, we focus on (i) extracting and (ii) disambiguating dataset mentions from social science publications to a list of given dataset references. Identifying dataset mention in literature is a challenging problem due to the huge number of styles of citing datasets. Although there are proposed standards for dataset citation in full-texts, researchers still ignore or neglect such standards (see, e.g., [13]). Furthermore, in many research publication, a correct citation of datasets is often missing [1]. The following two sentences exemplify the problem of the usage of an abbreviation to make a reference to an existing dataset. The first example illustrates the use of abbreviations that are known mainly in the author’s research domain. The latter illustrates the ambiguity of abbreviations. In this case, *WHO* identifies a dataset published by the World Health Organization and does not refer to the institution itself.

Example 1: *P-values are reported for the one-tail paired t-test on Allbus (dataset mention) and ISSP (dataset mention).*

Example 2: *We used WHO data from 2001 (dataset mention) to estimate the spreading degree of AIDS in Uganda.*

We treat the problem of detecting dataset mentions in full-text as a Named Entity Recognition (NER) task.

Formal problem definition Let D denote a set of existing datasets d and the knowledgebase K as a set of known dataset references k . Furthermore, each element of K is referencing an existing dataset d . The Named Entity Recognition and Linking task is defined as (i) the identification of dataset mentions m in a sentence, where m references a dataset d and (ii) linking them, when possible, to one element in K (i.e., the reference dataset list given by the RCC).

1.3.2 Challenges

We focus on the extraction of dataset mentions in the body of the full-text of scientific publications. There are three types of dataset mentions: (i) The full name of a dataset (“National Health and Nutrition Examination Survey”), (ii) an abbreviation (“NHANES”) or (iii) a vague reference, e.g., “the monthly statistic”. With all these types, the NER task faces special challenges. In the first case, the used dataset name can vary in different publications. For instance one publication cites the dataset with “National Health and Nutrition Examination Survey” the other could use the words “Health and Nutrition Survey”. In a case where abbreviations are used, a disambiguation problem occurs, e.g., in “WHO data”. WHO may describe the World Health Organization or the White House Office. The biggest challenge is again the lack of a gold standard that can be used to train a classifier. In the following we describe how we have dealt with this lack of ground truth data.

1.3.3 Phase one approach

Missing ground truth data is the main problem to handle during this competition. To this end, supervised learning methods for dataset mentions extraction from texts are not applicable without the identification of external training data or the creation

¹⁴ <https://github.com/CeON/CERMINE>

¹⁵ <https://jats.nlm.nih.gov>

of usefull labeled training data from information given by the competition. Because of the lack of existing training data for the task of dataset mention extraction we resort to the provided list of dataset mentions and publication pairs and re-annotate the particular sentences in the publication text. A list of identifying words is provided for some of the known links between publications and a datasets by the competition. These words represent the evidence of the linkage between publication and datasets and are extracted from the publication text. In the course of re-annotation, we search for each of the identifying words in the corresponding publication texts. For each match, we annotate the occurance in our raw text and use these annotations as ground truth. As described in the pre-processing section, our unit for processing the publication text are paragraphs. The re-annotated corpus consists of a list of paragraphs for each publication with stand-off annotations identifying the mentions of datasets (i.e. position of the start and end characters and the entity type for each mention: *dataset*). This re-annotation is then used to train Spacy’s neural network-based NER model¹⁶. We created a holdout set of 1,000 publications and a training set of size 4,000. Afterward, we train our model with the paragraphs as a sampling unit. In the training set, 0.45 percent of the paragraphs contained mentions. For each positive training example, we have added one negative sample that contains no known dataset mentions and is randomly selected. We used a batch size of 25 and a dropout rate of 0.4. The model was trained for 300 iterations.

Evaluation We evaluated our model with respect to four metrics: precision and recall, each for strict and for partial match. While the strict match metrics are standard evaluation metrics, the partial match metrics are their relaxed variants in which the degree to which dataset mentions have to match can vary. Consider the following partial match example: ”National Health and Nutrition Examination Survey” is the extracted dataset mention, while ”National Health and Nutrition Examination Survey (NHANES)” is the true dataset mention. In contrast to the strict version of the metrics, this overlapping match is considered a match for the partial version. The scores describe whether a model is able to find the correct positions of dataset mentions in the texts, even if the start and end positions of the characters are not the same, but the ranges overlap.

Table 2 show the results of the dataset mention extraction on the holdout set. The model can achieve high strict precision and recall values. As expected, the results are even better for the partial version of the metrics. It means that even if we couldn’t match the dataset mention in a text exactly, we can find the right context with very high precision.

1.3.4 Phase two approach

In the second phase of the competition, additional 5,000 publications were provided by RCC. We extended our approach to consider the list with dataset names supplied by the organizers and re-annotated the complete corpus of 15,000 publications) in the same manner as in phase one to obtain training data. This time we split the data in 80% for training and 20% for test.

Evaluation We resort to the same evaluation metrics as in phase one. However, we calculate precision and recall on the full-text of the publication and not on the paragraphs as in the first phase. Table 3 shows the results achieved by our model. We observe lower precision and recall values. Compared to phase one, there is also a smaller difference between the precision and recall values for the strict and partial version of the metrics.

1.4 Research method extraction

1.4.1 Task description

Inspired by a recent work of Nasar et al. [14], we define a list of basic entity types that give key-insights into scholarly publications. We adapted the list of semantic entity types to the domain of the social sciences with a focus on *research methods*, but also including related entity types such as *Theory*, *Model*, *Measurement*, *Tool*, *Performance*. We suspect that the division into semantic types might be helpful to find *research methods*. The reason is that the related semantic entities types might provide clues or might be directly related to the research method itself.

¹⁶ spacy.io

Table 2. Performance of phase one approach of dataset extraction.

Metric	Value
Precision (partial match)	0.93
Recall (strict match)	0.95
Precision (strict match)	0.80
Recall (strict match)	0.81

For example, in order to achieve a certain research goal, an experiment is used in which a certain combination of *methods* is applied to a *dataset*. The methods can be specified as concepts or indirectly through the use of certain *software*. The result is then quantified with a *performance* using a specific measure.

Example: *P-values* (measurement) are reported for the *one-tail paired t-test* (method) on *Allbus* (dataset) and *ISSP* (dataset). We selected the entity types *Research Method*, *Research Theory*, *Research Tool* and *Research Measurement* as the target research method related entity types (see Table 4). This decision is based on an examination of the SAGE ontology given by the RCC as a sample of how research method terms might look like.

Formal problem definition The task of Named Entity Recognition and Linking is to (i) identify the mentions *m* of research-related entities in a sentence and (ii) link them, if possible, to a reference knowledge base *K* (i.e. the SAGE Thesaurus¹⁷) or (iii) assign a type to each entity, e.g. a *research method*, selected from a set of predefined types.

1.4.2 Challenges

There are some major challenges that any named entity recognition, classification and linking system needs to handle. First, regarding NER, identifying the entities boundary is important, thus detecting the exact sequence span. Second, ambiguity errors might arise in classification. For instance, ‘range’ might be a domain-specific term from the knowledge base or belong to the general domain vocabulary. This is a challenging task for which context information is required. In the literature, this relates to the problem of **domain adaptation** which includes fine-tuning to specific named entity classes¹⁸. With respect to entity linking, another challenge is detecting name variations, since entities can be referred to in many different ways. Semantically similar words, synonyms or related words, which might be lexically or syntactically different, are often not listed in the knowledge base (e.g., the lack of certain terms like ‘questioning’ but not ‘questionnaire’). This problem of automatically detecting these relationships is generally known as **linking problem**. Note that part of this problem also results from PDF-to-text conversion which is error-prone. Dealing with incomplete knowledge bases, i.e. **handling of out of vocabulary (OOV) items**, is also a major issue, since knowledge bases are often not exhaustive enough and do not cover specific terms or novel concepts from recent research. Last but not least, the combination of different semantic types gives a more coherent picture of a research article. We hypothesize that such information would be helpful and results in an insightful co-occurrence statistics, and provides additional detail directly related to entity resolution, and finally helps to assess the **relevance of terms** by means of a score.

1.4.3 Our approach

Our research method extraction tool builds on Stanford’s CoreNLP and Named Entity Recognition System¹⁹. The information extraction process follows the workflow depicted in Figure 2, using separate modules for pre-processing, classification, linking and term filtering.

We envision the task of finding entities in scientific publications as a sequence labeling problem, where each input word is classified as being of a dedicated semantic type or not. In order to handle entities related to our domain, we train a CRF based machine learning classifier with major semantic classes, (see Table 4), using training material from the ACL RD-TEC 2.0 dataset [16]. Apart from this, we follow a domain adaptation approach inspired by [17] and ingest semantic background knowledge extracted from external scientific corpora, in particular the ACL Anthology [11, 18]. We perform entity linking by means of a new gazetteer based on a SAGE dictionary of Social Research Methods [19], thus putting a special emphasis on the social sciences. The linking component addresses the synonymy problem and matches an entity despite name variations such as spelling variations. Finally, term filtering is carried out based on termhood and unithood, while scoring is achieved by calculating a relevance score based on TF-IDF (cf. Section 1.4.3).

¹⁷ <http://methods.sagepub.com>

¹⁸ apart from those used in traditional NER systems like *Person*, *Location*, or *Organization* with abundant training data, as covered in the Stanford NER system[15]

¹⁹ <https://nlp.stanford.edu/projects/project-ner.shtml>

Table 3. Performance of phase two approach for dataset extraction.

Metric	Value
Precision (partial match)	0.51
Recall (partial match)	0.90
Precision (strict match)	0.49
Recall (strict match)	0.87

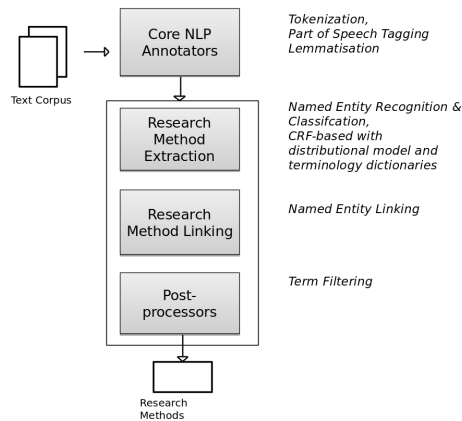


Figure 2. Overview of the entity extraction pipeline

Our research experiments are based on publications from the Social Science Open Access Repository (SSOAR)²⁰ as well as the train and test data of the Rich Context Competition corpus²¹. Our work extends previous work on this topic (cf. [20]) in various ways: First, we do not limit our study to abstracts, but use the entire fulltext. Second, we focus on a broader range of semantic classes, i.e. *Research Method*, *Research Theory*, *Research Tool* and *Research Measurement*, tackling also the problem of identifying novel entities.

Distributed semantic models For domain adaptation, we integrate further background knowledge. We use topical information from word embeddings trained on an scientific corpus as an additional features to our NER model. For this, we use agglomerative clustering of the word embeddings to identify topical groups of words. The cluster number of each word is used as additional sequential input feature for our CRF model. Semantic representations of words are a successful extension of common features, resulting in higher NER performance [21] and can be trained offline. In this work, the word vectors were learned based on 22,878 documents of the scientific ACL Anthology Reference Corpus²² using Gensim²³ with the skip gram model (cf. [22]) and a pre-clustering algorithm²⁴.

Features The features incorporated into the linear chain CRF are shown in the Table 5. The features depend mainly on the observations and on pairs of adjacent labels, using a log-linear combination. However, since simple token level training of CRFs leads to poor performance, more effective text features such as word shape, orthographic, gazetteer, Part-Of-Speech (POS) tags, along with word clustering (see Section 1.4.3) have been used.

Knowledge resources We use the SAGE thesaurus which includes well-defined concepts, an explicit taxonomic hierarchy between concepts as well as labels that specify synonyms of the same concept. A portion of terms is unique to the social science domain (e. g., ‘dependent interviewing’), while others are drawn from related disciplines such as statistics (e. g., ‘conditional

²⁰ <https://www.ssoar.info>

²¹ <https://coleridgeinitiative.org/richcontextcompetition> with a total of 5,000 English documents

²² <https://acl-arc.comp.nus.edu.sg/>

²³ <https://radimrehurek.com/gensim/>

²⁴ Word embeddings are trained with a skip gram model using embedding size equal to 100, word window equal to 5, minimal occurrences of a word to be considered 10. Word embeddings are clustered using agglomerative clustering with a number of clusters set to 500, 600, 700. Ward linkage with Euclidean distance is used to minimize the variance within the clusters.

Table 4. Entity types of relevance for the research method extraction task.

Entity type	Corresponding SAGE type	Incl. statistic glossary	Examples
Research Method	SAGE-METHOD	✓	Bootstrapping, Active Interviews
Research Measurement	SAGE-MEASURE		Latent Variables, Phi coefficient, Z-score
Research Theory	SAGE-THEORY		Frankfurt shool, Feminism, Actor network theory
Research Tool	SAGE-TOOL		SPSS, R statistical package

Table 5. Features used for NER

Type	Features
Token unigrams	$w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2}, \dots$
POS unigrams	p_i, p_{i-1}, p_{i-2}
Shapes	shape and capitalization
NE-Tag	t_{i-1}, t_{i-2}
WordPair	(p_i, w_i, c_i)
WordTag	(w_i, c_i)
Gazetteer	SAGE Gazetteer
Distributional Model	ACL Anthology model

Table 6. Most relevant terms from SAGE by Semantic Type

SAGE Term	TF-IDF Score	Semantic Class
Fuzzy logic	591,29	Research Method
arts-based research	547,21	Research Method
cognitive interviewing	521,13	Research Method
QCA	463,13	Research Method
oral history	399,68	Research Method
market research	345,37	Research Field
life events	186,61	Research Field
Realism	314,34	Research Theory
Marxism	206,77	Research Theory
ATLAS.ti	544,51	Research Tool
GIS	486,01	Research Tool
SPSS	136,52	Research Tool

likelihood ratio test’²⁵. However, since the thesaurus is not exhaustive and covers only the top-level concepts related to social science methods, our aim was to extend it by automatically extracting further terms from domain-specific texts, in particular from the Social Science Open Access Repository. More concretely, we carried out the following steps to extend SAGE as an off-line step. For step 2 and 3, candidate terms have been extracted by our pipeline for the entire SSOAR corpus.

1. Assignment of semantic types to concepts (manual)
2. Extracting terms variants such as abbreviations, synonyms, related terms from SSOAR (semi-automatic)
3. Computation of term and document frequency scores for SSOAR (automatic)

Extracting term variants such as abbreviations, synonyms, and related terms 26,082 candidate terms have been recognized and classified by our pipeline and manually inspected to a) find synonyms and related words that could be linked to SAGE, and b) build a post-filter for incorrectly classified terms. Moreover, abbreviations have been extracted using the algorithm of Schwartz and Hearst [23]. This way, a Named Entity gazetteer could be built and is used at run-time. It comprises 1,111 terms from SAGE and 447 terms from the used glossary of statistical terms²⁶ as well as 54 previously unseen terms detected by the model-based classifier.

Computation of term and document frequency scores Term frequency statistics have been calculated off-line for the entire SSOAR corpus. The term frequency at corpus level will be used at run time to determine the term relevance at the document level by calculating the TF-IDF scores. The most relevant terms from SAGE are listed in Table 6.

Definition of a relevance score Relevance of terminology is often assessed using the notion of *unithood*, i.e. ‘the degree of strength or stability of syntagmatic combinations of collections’, and *termhood*, i.e. ‘the degree that a linguistic unit is related to domain-specific concepts’ [24]. Regarding *unithood*, the NER model implicitly contains heuristics about legal POS tag

²⁵ A glossary of statistical terms as provided in <https://www.statistics.com/resources/glossary/> has been added as well.

²⁶ Based on <https://www.statistics.com/resources/glossary>

sequences for candidate terms, consisting of at least one noun (NN), preceded or followed by modifiers such as adjectives (JJ), participles (VB*) or cardinal numbers (CD), complemented by wordshape features.

In order to find out if the candidate term also fulfills the *termhood* requirement, domain-specific term frequency statistics have been computed on the SSOAR repository, and set in contrast to general domain vocabulary terms. It has to be noted that only a small portion of the social science terms is actually unique to the domain (e.g., ‘dependent interviewing’), while others might be drawn from related disciplines such as statistics (e.g., ‘conditional likelihood ratio test’).

Preliminary results Our method has been tested on 100 fulltext papers from SSOAR and ten documents from the Rich Context Competition (RCC), all randomly selected from a hold out corpus. In our experiments on SSOAR Social Science publications, we compared results to the given metadata information. The main finding was that while most entities from the SAGE thesaurus could be extracted and linked reliably (e.g., ‘Paired t-test’), they could not be easily mapped to the SSOAR metadata terms, which consist of only a few abstract classes (e.g., ‘quantitative analysis’). Furthermore, our tool was tested by the RCC organizer, where the judges reviewed ten random publications and generated qualitative scores for each document.

1.5 Research field classification

1.5.1 Task description

The goal of this task is to identify the research fields covered in the social science publications. In general, two approaches could be applied to this task. One is the extraction of relevant terms of the publications. It means that this task could be seen as a keyword extraction task and the detected terms considered as descriptive terms regarding the research field. The second approach is to learn to classify publications research fields with the use of annotated data in a supervised manner. The benefit of the second approach is that the classification scheme to describe the research field can be defined by experts of the domain. The disadvantage of supervised trained classifiers for this task is the lack of applicable training data. Furthermore, it must be ensured that the training data is comparable to the texts the research field classifier should be applied on.

Formal problem definition Let P denote a set of publications of size n , A a set of corresponding abstracts of the same size and L a set of k defined class labels describing research fields. The task of research field classification is to select for each publication $p_i \in P$ based on the information contained in the corresponding abstract $a_i \in A$ a set of labels $C_i = \emptyset \cap \{c_1 \dots c_n | c_n \in L\}$ of n labels. The number of n denotes the number of labels from L describing the research field of a_i and can vary for each publication p_i . If there is no label l_k representing the information given by the abstract a_i the set of class labels is the empty set \emptyset .

1.5.2 Our approach

Since we didn’t receive any gold standard for this task during the competition we decided to make use of external resources. We decided to use an external labeled dataset to train a text classifier which is able to predict one or more research label for a given abstract of a publication.

The publications given throughout the competition belongs to the domain of social sciences we considered metadata from an open access repository for the social sciences called SSOAR. The advantages are twofold. On the one hand, we could rely on professional annotations in a given classification scheme covering the social sciences and related areas. On the other hand the source is openly available.²⁷

The annotated data of SSOAR contains four different annotation schemes for research field related information. By reviewing these schemes, we decided to use the Classification Social Science (classoz) annotation scheme. The number of classes in each schema, coverage of each classification, and the distribution of data in each schema affected our decision. An exhaustive description of the used data can be found in 1.2.2.

Pre-processing and model architecture SSOAR is a multilingual repository. Therefore, the available abstracts may vary in language and the language of the abstract may differ from the language of the article itself. We selected all English abstract with valid classification as our dataset. Mainly because of the language of the RCC corpus. However, it should be noted that the multilingual SSOAR abstract corpus has a skewed distribution of languages with English and German as the main languages. We count 22,453 English abstracts with valid classification after filtering. Due to the unequal distribution of labels in the dataset, we need to guaranty enough training data for each label. We selected only labels with frequency over 300 for training the model, which results in a total of 44 out of 154 classification labels representing research fields. For creating train and test set, 22,453 SSOAR publications with their assigned labels were split randomly. We used a train/validation/test split of 70/10/20. We decided to train a text classifier based on a fasttext [10] model in the author’s implementation. The arguments to use this model was the speed in comparison to a more complex neural net architecture and the still comparable to state of the art performance (e.g.[25]). The model is trained with learning rate 1.0 for 150 epochs. Also, the negative sampling parameter is set to 25.

²⁷ A script to download the metadata of SSOAR can be found in [github/research-field-classifier](https://github.com/research-field-classifier)

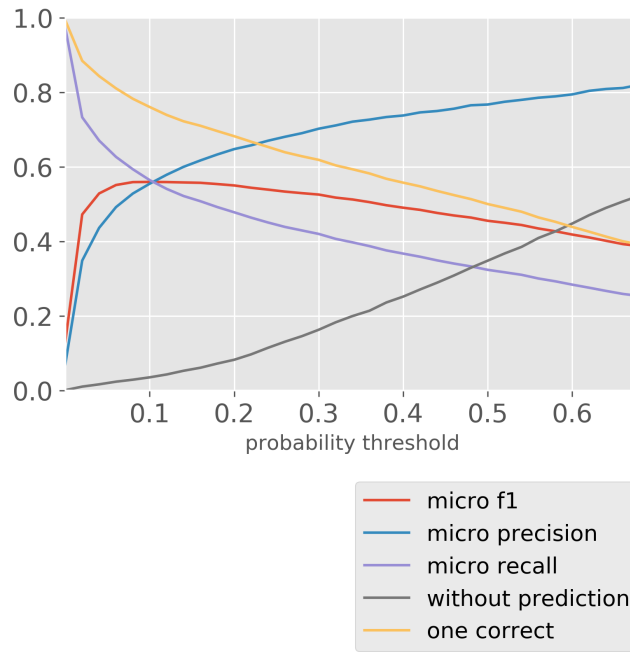


Figure 3. Metrics for different selected probability thresholds (validation set)

1.5.3 Evaluation

Figure 3 shows the performance of the model regarding various evaluation metrics for different thresholds. A label is assigned to a publication if the model outputs a probability for the label above the defined threshold. In multi-label classification, this allows us to evaluate our model from different perspectives. As illustrated in figure 3, the intersection of the micro precision and the micro recall curves is at the threshold of 0.1, where the highest micro f1 score is achieved. By increasing the threshold from this point, the micro-precision score is increasing, but the micro recall is falling. By decreasing threshold, these trends are inverted. Also, the default threshold of 0.5 doesn't look promising. In spite of micro-precision about 0.75, we have a problem with the very high number of items without any prediction. In respect to this observation it is advantageous to select a lower threshold in a productive environment. The curve named *without prediction* shows for a given threshold the share of publications in the test set without any prediction. If the selected threshold value is high, the number of publications for which the model cannot predict a research field increases. For example, a selected threshold value of 0.55 leads to 40% unclassified publications in the test set. The *one correct* named curve indicates the quality of the publication wise prediction. It shows the share of all publications in the test set where at least one of the predicted research field labels can be found in the ground truth data. For instance, if a threshold of 0.1 is selected for 75% of the publications in the test set, at least one of the model predictions are correct. This value decreases with increasing threshold similar to the recall metric. The final metrics for our model and a selected threshold of 0.1 is shown in table 7. The final micro f1 value on the test set is 0.56.

1.6 Discussion

1.6.1 Dataset Extraction

For the dataset extraction task the proposed method are only tested on social science related data. The performance measures we have introduced are based on a hold out data set of our automatically created dataset. Especially the recall could be biased.

Table 7. Evaluation results on test set

Metric	Value
Micro precisison:	0.554
Micro recall:	0.564
Micro f1:	0.559

This is because we only label known data set. The results of the second phase presented during the RCC workshop²⁸ are showing good performance of our approach in comparison to the other finalist teams with the highest precision 52.2% (second: 47.0%) and second in recall (ours: 20.5, best: 34.8%). This lead to the second best performing system for this task in respect of f1 measure (29.5%, 40.0% first place). The results on the manually created hold out set pointing out, that our system performs better in respect to precision in comparison to the other finalist teams. The social science focused corpus of research publications and dataset metadata lead us to suppose, that our trained model is working bad on different domains of science. Especially the focus on survey data and the reflection in dataset names (e.g. Current Population Survey) could have biased the model to detect the survey as a specific type of research datasets better than other subtypes like e.g. text corpora in the NLP community. In general our approach to use a weakly labeled corpus created using a list of dataset names could be applied in other research domains. The needed resources are a set of full text publications and a sufficiently large list of dataset names for this domain.

1.6.2 Research method extraction

The extraction of research methods from full text publications we consider as the most challenging task. This is because, the sample vocabulary given by the competition covers a large thematic area, from dataset, over mathematical models to qualitative methods. The task itself was defined as the identification of research methods associated to the publication. On the one hand we were confronted with a lack of training data in the competition. On the other hand, in contrast to the task of research field classification, we were not able to identify external corpora which could be applied on this task. We combined models from another research domain with a manually curated extension of known research method terms. The qualitative reviews during the two phases of the competition attested that this approach works. A valid quantitative evaluation is prevented by the lack of ground truth data.

1.6.3 Research field classification

Our supervised machine learning approach to handle the research field classification task performs well on the dataset created from social science publication metadata. A micro f1 measure of above 55% seems to be a good result for a dataset with 44 labels and a mean number of keywords of three terms per publication. As one example of multilabel classification with a comparable size of labels we would like to mention the classification of texts in the domain of medicine presented in [25]. The models tested by the authors on the task of multilabel prediction from 50 different labels leads to micro f1 values between 53% and 62%.

The analysis of the performance of our model does not enable us to determine the performance on other research domains than the social sciences. Even if the used classification scheme covers neighbour disciplines like medicine, the numbers of samples of the training data covering other research fields than the social science is less. If one considers this fact, it can be assumed that the performance for corpora of other disciplines is lower. A benefit of our approach is the usage of abstracts as input to classify the research publications. This make the approach usefull even if there are no full text of publications are available. On the other hand, the use of Cermine to extract information from publications available as PDF files enables us to automatically extract abstracts. With the help of this we are able to classify publications, even if abstracts not present in the publications metadata.

1.7 Conclusion

1.8 Acknowledgments

We would like to thank GESIS for giving us the time and resources to participate in the competition. The work share of the first author was funded by Deutsche Forschungsgemeinschaft (DFG) under grant number MA 3964/7-1.

1.9 References

References

- [1] Katarina Boland, Dominique Ritze, Kai Eckert, and Brigitte Mathiak. Identifying references to datasets in publications. In *International Conference on Theory and Practice of Digital Libraries*, pages 150–161. Springer, 2012.
- [2] Behnam Ghavimi, Philipp Mayr, Christoph Lange, Sahar Vahdati, and Sören Auer. A semi-automatic approach for detecting dataset references in social science texts. *Information Services & Use*, 36(3-4):171–187, 2016.
- [3] Tobias Backes. The impact of name-matching and blocking on author disambiguation. In Alfredo Cuzzocrea, James Allan, Norman W. Paton, Divesh Srivastava, Rakesh Agrawal, Andrei Z. Broder, Mohammed J. Zaki, K. Selçuk Candan, Alexandros Labrinidis, Assaf Schuster, and Haixun Wang, editors, *CIKM*, pages 803–812. ACM, 2018.

²⁸ Agenda of the Workshop: <https://coleridgeinitiative.org/richcontextcompetition/workshopagenda>. The results of the finalists are presented here: <https://youtu.be/PE3nFrEkwoU?t=9865>.

- [4] Tobias Backes. Effective unsupervised author disambiguation with relative frequencies. In Jiangping Chen, Marcos André Gonçalves, Jeff M. Allen, Edward A. Fox, Min-Yen Kan, and Vivien Petras, editors, *JCDL*, pages 203–212. ACM, 2018.
- [5] Katarina Boland and Frank Krüger. Distant supervision for silver label generation of software mentions in social scientific publications. In *Proceedings of the 4th Joint Workshop on Bibliometric-enhanced Information Retrieval and Natural Language Processing for Digital Libraries*, pages 15–27, 2019.
- [6] Ran Yu, Ujwal Gadiraju, Besnik Fetahu, Oliver Lehmberg, Dominique Ritze, and Stefan Dietze. Knowmore - knowledge base augmentation with structured web markup. *Semantic Web*, 10(1):159–180, 2019.
- [7] Pracheta Sahoo, Ujwal Gadiraju, Ran Yu, Sriparna Saha, and Stefan Dietze. Analysing structured scholarly data embedded in web pages. *Lecture Notes in Computer Science*, 9792, 05/2017 2017.
- [8] Daniel Hienert, Dagmar Kern, Katarina Boland, Benjamin Zapilko, and Peter Mutschke. A digital library for research data and related information in the social sciences. In Maria Bonn, Dan Wu, J. Stephen Downie, and Alain Martaus, editors, *JCDL*, pages 148–157. IEEE, 2019.
- [9] Thomas Krämer, Claus-Peter Klas, and Brigitte Hausstein. A data discovery index for the social sciences. In *Scientific data*, 2018.
- [10] Armand Joulin, Edouard Grave, Piotr Bojanowski, and Tomas Mikolov. Bag of tricks for efficient text classification. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*, pages 427–431. Association for Computational Linguistics, April 2017.
- [11] Steven Bird, Robert Dale, Bonnie J Dorr, Bryan Gibson, Mark Thomas Joseph, Min-Yen Kan, Dongwon Lee, Brett Powley, Dragomir R Radev, and Yee Fan Tan. The acl anthology reference corpus: A reference dataset for bibliographic research in computational linguistics. In *Proceedings of the Sixth International Conference on Language Resources and Evaluation (LREC 2008)*. European Language Resources Association (ELRA), 2008.
- [12] Dominika Tkaczyk, Paweł Szostek, Mateusz Fedoryszak, Piotr Jan Dendek, and Łukasz Bolikowski. Cermine: automatic extraction of structured metadata from scientific literature. *International Journal on Document Analysis and Recognition (IJDAR)*, 18(4):317–335, 2015.
- [13] Micah Altman and Gary King. A proposed standard for the scholarly citation of quantitative data. *D-lib Magazine*, 13(3/4), 2007.
- [14] Zara Nasar, Syed Waqar Jaffry, and Muhammad Kamran Malik. Information extraction from scientific articles: a survey. *Scientometrics*, 117(3):1931–1990, 2018.
- [15] Jenny Rose Finkel, Trond Grenager, and Christopher Manning. Incorporating non-local information into information extraction systems by gibbs sampling. In *Proceedings of the 43rd annual meeting on association for computational linguistics*, pages 363–370. Association for Computational Linguistics, 2005.
- [16] Behrang QasemiZadeh and Anne-Kathrin Schumann. The acl rd-tec 2.0: A language resource for evaluating term extraction and entity recognition methods. In *LREC*, 2016.
- [17] Rodrigo Agerri and German Rigau. Robust multilingual named entity recognition with shallow semi-supervised features. *Artificial Intelligence*, 238:63–82, 2016.
- [18] Daniel Gildea, Min-Yen Kan, Nitin Madnani, Christoph Teichmann, and Martin Villalba. The acl anthology: Current state and future directions. In *Proceedings of Workshop for NLP Open Source Software (NLP-OSS)*, pages 23–28, 2018.
- [19] Michael Lewis-Beck, Alan E Bryman, and Tim Futing Liao. *The Sage encyclopedia of social science research methods*. Sage Publications, 2003.
- [20] Judith Eckle-Kohler, Tri-Duc Nghiem, and Iryna Gurevych. Automatically assigning research methods to journal articles in the domain of social sciences. In *Proceedings of the 76th ASIS&T Annual Meeting: Beyond the Cloud: Rethinking Information Boundaries*, page 44. American Society for Information Science, 2013.
- [21] Joseph Turian, Lev Ratinov, and Yoshua Bengio. Word representations: A simple and general method for semi-supervised learning. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics, ACL '10*, pages 384–394, Stroudsburg, PA, USA, 2010. Association for Computational Linguistics.
- [22] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119, 2013.
- [23] Ariel S. Schwartz and Marti A. Hearst. A simple algorithm for identifying abbreviation definitions in biomedical text. In *Pacific Symposium on Biocomputing*, pages 451–462, 2003.

- [24] Kyo Kageura and Bin Umino. Methods of automatic term recognition: A review. *Terminology. International Journal of Theoretical and Applied Issues in Specialized Communication*, 3(2):259–289, 1996.
- [25] Guoyin Wang, Chunyuan Li, Wenlin Wang, Yizhe Zhang, Dinghan Shen, Xinyuan Zhang, Ricardo Henao, and Lawrence Carin. Joint embedding of words and labels for text classification. *arXiv preprint arXiv:1805.04174*, 2018.